
From Circuits to Claims: Post-Selection Certificates for Mechanistic Interpretability

Anonymous Author(s)

Affiliation

Address

email

Abstract

Circuit discovery in mechanistic interpretability increasingly relies on adaptive selectors, counterfactual datasets, and standardized evaluations. We study selected circuits as post-selection objects and define a mechanistic claim by its model, computational graph, task distribution, intervention design, task score, selector, selection data, selected circuit, and audited properties. We introduce Circuit Certificates, a claim-centered audit framework reporting behavioral preservation, localization, edgewise minimality, minimality profiles, selector stability, intervention robustness, power, and structural recovery when ground truth is available.

The framework implements fixed holdout, folded fixed-circuit audit, selective split, and selector-level cross-fit protocols. Selective split yields finite-sample post-selection validity for any measurable selector under independent audit data. Selector-level cross-fit audits the selection procedure through a procedure-level risk estimator. Edgewise minimality is assessed with Holm-adjusted tests, with minimality profiles summarizing the distribution of edge effects.

We instantiate the framework in an audited artifact covering 65 run packets and 158 certificates. The results identify distinct evidence regimes across tasks. On GPT-2 small IOI, selector-level cross-fit certifies localization with high selector stability and adequate power, while preservation and minimality are uncertified. On Gemma2 ARC-Easy, it certifies preservation while localization is uncertified. On Gemma2 arithmetic subtraction, it certifies localization in a lower-power regime. On InterpBench, behavioral support coincides with weak edge-level structural recovery. These findings support reporting mechanistic interpretability results as fully specified audited claims.

1 Introduction

Circuit analysis has become a central mode of mechanistic interpretability. Language-model circuits are discovered with activation patching, path patching, attribution-based surrogates, and mask-optimization methods. Canonical tasks such as indirect object identification and arithmetic now have established datasets and counterfactual constructions. Benchmark suites increasingly standardize tasks, metrics, and selectors [1, 4–6, 12, 14, 15]. This infrastructure makes circuit discovery more systematic and sharpens the evidential question attached to any returned subgraph: which mechanistic claim does it support?

A selected circuit is the output of an adaptive pipeline whose construction depends on graph granularity, task distribution, counterfactual construction, scoring rule, selector class, calibration rule, and often a size threshold. The returned subgraph therefore inherits the assumptions and degrees of

freedom introduced by those choices, while its evidential status depends on the audit protocol, the relation between audit and selection data, the intervention family used for evaluation, the multiplicity induced by edgewise testing, the selector’s stability under resampling, and, when available, the relation between behavioral support and structural recovery. Recent work identifies these axes as sources of inferential variation: intervention design changes the estimand of faithfulness metrics [7, 8, 11]; candidate-circuit evaluation admits a formal testing treatment [14]; circuit discovery behaves as a statistical estimation pipeline with substantial instability [10]; and benchmarks separate causal localization from structural recovery [5, 12]. A reporting standard for circuit discovery should specify the claim, the selection procedure, the audit target, and the uncertainty attached to the result.

The framework introduced here takes the *mechanistic claim* as its central object: a selected circuit together with its model, graph abstraction, task distribution, intervention design, score, selector, selection data, and audited property. Specifying claims at this level gives post-selection inference a direct role in circuit auditing. The resulting framework, which we call *Circuit Certificates*, addresses the following research question:

Given a discovered circuit, which claim does the audit support, under which intervention design and protocol, and with what statistical uncertainty?

Formalization. We define a mechanistic claim as an inferential object indexed by a model, graph, task distribution, intervention design, task score, selector, selection data, and selected circuit. This vocabulary assigns separate audit targets to behavioral preservation, localization, edgewise minimality, selector stability, intervention robustness, power, and structural recovery.

Theory. We develop four audit protocols: fixed holdout, folded fixed-circuit audit, selective split, and true selector-level cross-fit. The audit target may be a fixed circuit or the selector that produces circuits. For selective split, we prove finite-sample validity under independent audit data for any measurable selector. For true selector-level cross-fit, we identify the target as a procedure-level out-of-fold estimand and separate foldwise validity from stronger aggregate concentration claims. We also obtain familywise-valid edgewise minimality decisions through Holm adjustment.

Evidence. We instantiate the framework in a protocol-audited artifact and evaluate it on canonical and benchmarked settings. The empirical results show repeated separations between evidence dimensions that are often compressed into a single circuit-quality score.

The audits yield distinct claim profiles across regimes. On GPT-2 small IOI, true selector-level cross-fit supports localization with high selector stability and adequate power; preservation, minimality, and strong robustness remain unsupported. On Gemma2 ARC-Easy, the audit supports preservation with adequate power and leaves localization unsupported. On Gemma2 arithmetic subtraction, a second non-IOI case exhibits selector-level localization in a weaker, low-resolution regime. On InterpBench, selector-level behavioral support coexists with weak edge-level structural recovery, separating behavioral support from ground-truth circuit recovery. A final MCQA control shows that selector consistency contributes one layer of evidence within a broader certificate. Together, these results support a reporting practice in which discovered circuits are presented through audited claim profiles.

The framework builds on the circuit perspective of Elhage et al. [2], Olah et al. [13]; the intervention-aware methodology of path patching and activation patching [4, 7]; the semantic foundations supplied by causal abstraction [3]; fixed-circuit testing in Shi et al. [14]; and the empirical anchors supplied by InterpBench and MIB [5, 12]. The result is a rigorous and executable framework for moving from discovered circuits to auditable mechanistic claims.

2 Mechanistic claims and audit protocols

2.1 Claims

Let M be a model represented at a fixed graph granularity by a directed acyclic graph $\mathcal{G} = (V, \mathcal{E})$. Let $\mathcal{D}$ denote a task distribution over tuples $(x, \tilde{x}, y)$, where x is a clean input, $\tilde{x}$ is a counterfactual input used for intervention, and y is task-level ground-truth information when required by the score. Let s be a scalar task score, and let η denote an intervention design, such as complement patching with

counterfactual activations, mean ablation, or optimal ablation. A selector $\mathcal{A}$ maps selector-visible data to a circuit $C \subseteq \mathcal{E}$.

Definition 1 (Mechanistic claim). *A mechanistic claim is a tuple*

$$\Gamma = (M, \mathcal{G}, \mathcal{D}, \eta, s, \mathcal{A}, S, C, \mathcal{P}),$$

where S is the selector-visible data, $C = \mathcal{A}(S)$ is the selected circuit, and $\mathcal{P}$ is a family of audited properties.

The tuple records the information needed to interpret the evidence attached to the selected circuit, including the model, graph abstraction, task distribution, intervention design, score, selector, selection data, and audited properties. Because graph choices, counterfactual constructions, scores, and intervention families determine what the selected subgraph is evidence for, the audit treats the claim as the evidential unit.

2.2 Behavioral properties

We audit two behavioral properties that are central to circuit evaluation. Let $M_C^\eta(x, \tilde{x})$ denote the output of the selected circuit under intervention design η , and let $M_{\bar{C}}^\eta(x, \tilde{x})$ denote the output of the complement construction under the same design. Let ℓ be a bounded discrepancy between scores, logits, or task outputs.

Behavioral preservation. We define

$$R_{\text{pres}}^\eta(C) = \mathbb{E}_{(x, \tilde{x}, y) \sim \mathcal{D}} [\ell(M(x), M_C^\eta(x, \tilde{x}))].$$

Small values indicate that the selected circuit preserves the behavior of the full model under the chosen intervention design. This quantity measures behavioral sufficiency for the selected circuit.

Localization. We define

$$R_{\text{loc}}^\eta(C) = \mathbb{E}_{(x, \tilde{x}, y) \sim \mathcal{D}} [\ell(M(x), M_{\bar{C}}^\eta(x, \tilde{x}))].$$

Large values indicate that the complement construction loses the target behavior, so the behavior is concentrated in the selected circuit under the chosen intervention design.

Preservation quantifies behavioral sufficiency. Localization quantifies behavioral concentration. A circuit may preserve the target behavior while its complement also preserves it; another may concentrate the target behavior while failing to reproduce the full model on its own. The experiments below exhibit both regimes.

Edgewise minimality. For each edge $e \in C$, define the edge effect

$$\Delta_e(C; \eta) = R_{\text{pres}}^\eta(C \setminus \{e\}) - R_{\text{pres}}^\eta(C).$$

This quantity measures the loss in behavioral preservation produced by removing edge e from the selected circuit. Large circuits induce a family of edgewise tests, so minimality auditing requires multiplicity control.

Minimality profiles. A scalar minimality decision gives a coarse summary of edgewise structure. We report the profile

$$\text{MinProf}_C(\epsilon) = \frac{1}{|C|} \sum_{e \in C} \mathbf{1}[\Delta_e(C; \eta) > \epsilon].$$

As ϵ increases, $\text{MinProf}_C(\epsilon)$ records the fraction of edges whose effect exceeds the threshold. The profile distinguishes circuits with a descriptive high-effect tail from circuits with a broader low-effect fringe.

Selector stability. Given bootstrap or foldwise selector outputs $C^{(1)}, \dots, C^{(B)}$, stability is summarized by inclusion probabilities $\pi_e = \frac{1}{B} \sum_b \mathbf{1}[e \in C^{(b)}]$, mean pairwise Jaccard overlap, selection entropy, and thresholded core/fringe edge sets. These summaries provide a circuit-valued analogue of stability selection [9].




Figure 1: Audit protocols.

Intervention robustness. Intervention design changes the estimand [7, 8, 11], so claims are indexed by an intervention family $\mathcal{H}$. A certificate reports preservation and localization under each member of $\mathcal{H}$, together with preservation robustness, localization robustness, strong robustness, and a sensitivity norm.

Structural recovery. When a benchmark provides a ground-truth circuit C^* , the selected circuit can also be evaluated structurally through precision, recall, and F1. Structural recovery is reported alongside behavioral auditing, since edge overlap and behavioral support measure distinct properties of the selected circuit.

2.3 Audit protocols

Each claim includes an audit protocol, chosen from four cases:

1. **Fixed holdout:** a fixed externally supplied circuit is audited on independent holdout data.
2. **Folded fixed-circuit audit:** a fixed circuit is audited over multiple folds, with selection performed once.
3. **Selective split:** discovery and calibration are performed on selector-visible data, and the selected circuit is audited on an independent audit split.
4. **True selector-level cross-fit:** for each fold, the circuit is rediscovered without using that fold and audited on the held-out fold.

Selective split supports a selected-circuit claim for the circuit returned by one adaptive pipeline. True selector-level cross-fit audits the selection procedure by repeatedly rediscovering circuits and auditing them on held-out folds. The certificate records the protocol so that each reported property carries the corresponding evidential scope.

3 Circuit Certificates

A *Circuit Certificate* is the audited evidence profile attached to a mechanistic claim. We use *certificate* in the statistical-audit sense: a structured, protocol-indexed record of the statements supported by the audit for the specified claim. The certificate records point estimates, one-sided confidence bounds, protocol metadata, minimality outputs, stability summaries, robustness vectors, structural-alignment results when available, and power calculations.

150 3.1 Selective-split validity

151 Selective split is the basic post-selection protocol. Discovery and calibration may be adaptive; the
 152 audit data are held out from both stages. This separation yields finite-sample one-sided certification
 153 for any measurable selector.

154 **Theorem 2** (Selective-split post-selection validity). *Let $S = (D_{\text{disc}}, D_{\text{cal}})$ be selector-visible data
 155 and let $C = g(S)$ be any measurable selector output. Let D_{aud} be independent of S . For any
 156 bounded loss $w(C, Z) \in [a, b]$, define the audit risk*

$$R(C) = \mathbb{E}_{Z \sim \mathcal{D}}[w(C, Z)].$$

157 *Conditional on S , the audit losses are i.i.d. bounded observations with mean $R(C)$. Therefore any
 158 valid one-sided concentration bound for bounded i.i.d. means yields a finite-sample valid confidence
 159 bound for the selected circuit. In particular, empirical-Bernstein upper and lower confidence bounds
 160 are valid for preservation and localization, respectively.*

161 The theorem applies to arbitrary measurable selection pipelines, including manual, attribution-based,
 162 graph-pruning, and calibrated threshold procedures. Its identifying condition is independence between
 163 the audit data and the data used for discovery and calibration.

164 3.2 Selector-level cross-fit

165 True selector-level cross-fit audits the selector as a procedure. For folds $\mathcal{I}_1, \dots, \mathcal{I}_K$, let C_k denote
 166 the circuit selected without using fold k , and let $w(C_k, Z_i)$ be the corresponding audit loss for $i \in \mathcal{I}_k$.
 167 The procedure-level target is

$$R_{\text{cf}} = \frac{1}{n} \sum_{k=1}^K \sum_{i \in \mathcal{I}_k} \mathbb{E}[w(C_k, Z_i) \mid C_k].$$

168 The cross-fit estimator is

$$\hat{R}_{\text{cf}} = \frac{1}{n} \sum_{k=1}^K \sum_{i \in \mathcal{I}_k} w(C_k, Z_i).$$

169 **Proposition 3** (Foldwise out-of-fold validity). *For each fold k , conditional on the foldwise selected
 170 circuit C_k , the held-out losses $\{w(C_k, Z_i) : i \in \mathcal{I}_k\}$ are i.i.d. bounded variables. Consequently, one-
 171 sided empirical-Bernstein bounds can be applied validly on each held-out fold for its corresponding
 172 foldwise risk.*

173 Fold-resampled auditing provides selector-level evidence by estimating the behavior of the selection
 174 procedure under repeated foldwise discovery and held-out auditing. For the fully aggregated selector-
 175 level estimator, an unconditional finite-sample concentration result requires additional dependence-
 176 aware assumptions, since data from one audit fold may contribute to the selection of a circuit audited
 177 on another fold. The results in Section 5 use this distinction to separate fixed-circuit claims from
 178 procedure-level claims.

179 3.3 Edgewise minimality and multiplicity

180 Minimality is audited through edgewise hypotheses

$$H_{0,e} : \Delta_e(C; \eta) \leq \tau_{\min}, \quad H_{1,e} : \Delta_e(C; \eta) > \tau_{\min}, \quad e \in C.$$

181 The framework reports both Holm-adjusted and Benjamini–Hochberg summaries. The formal
 182 guarantee is familywise error control.

183 **Proposition 4** (Familywise-valid minimality decisions). *Under selective split or true selector-level
 184 cross-fit, if each edgewise p -value is valid conditional on the selected circuit and the audit sample
 185 used for that edge, then Holm adjustment controls the familywise error rate at level α for the family
 186 of edgewise minimality claims.*

187 A familywise edgewise decision yields a strong fixed-circuit certificate, while minimality profiles
 188 provide a complementary descriptive view of how edge effects are distributed across the selected
 189 circuit. Under selector-level cross-fit, edgewise minimality is interpreted foldwise or on a separately
 190 audited final circuit.

Case	n_{audit}	$(\tau_{\text{pres}}, \tau_{\text{loc}})$	$[a, b]$	Interventions
IOI	1000	(0.3, 0.3)	$[-2, 2]$	{cf, mean}
ARC-Easy	570	(0.2, 0.2)	$[-2, 2]$	{cf, mean}
Arithmetic	137	(0.2, 0.2)	$[-2, 2]$	{cf, mean}
InterpBench	1000	(0.2, 0.2)	$[-2, 2]$	{cf, mean}

Table 1: Main selector-level audit design. The backend score is `prob_diff` in $[-1, 1]$, and the bounded audited loss is the signed transformed score `prob_diff_delta` in $[-2, 2]$. Splits are formed at the unit of independent task-instance rows.

3.4 Power

Certificates also report the inferential resolution of the audit. For a target effect size δ , the framework computes minimum detectable effects under normal, Hoeffding, and empirical-Bernstein rules, with the selected rule recorded in the certificate. The main experiments use the empirical-Bernstein rule, aligning the power calculation with the confidence intervals used for preservation and localization.

Interpretation. A certificate reports mechanistic evidence as a profile over preservation, localization, minimality, stability, robustness, power, and structural recovery. These dimensions can agree, separate, or conflict. The experiments show how profile-level reporting states mechanistic claims with explicit scope and uncertainty.

4 Experimental design

Implementation. The framework is instantiated in a consolidated artifact built on TransformerLens-based discovery integrations and protocol-aware certificate builders, with explicit protocol metadata, claim and certificate schemas, selector integrations, task loaders, and archived run packets.

Selectors and interventions. The main experiments use EAP-IG as the primary selector and EAP as a comparator on InterpBench, following current practice in scalable circuit localization [6, 12]. The primary intervention is counterfactual patching, with mean ablation and optimal ablation used as robustness follow-ups.

Tasks. We evaluate four representative settings. **IOI** on GPT-2 small is the flagship canonical case [15]. **ARC-Easy** on Gemma2 provides a realistic task with strong selector-level preservation. **Arithmetic subtraction** on Gemma2 provides a second non-IOI localization case. **InterpBench IOI** provides semi-synthetic transformers with ground-truth circuits, enabling direct comparison between behavioral support and structural recovery [5]. A final MCQA control is reported in the appendix.

Protocols. The main table reports true selector-level cross-fit, the primary selector-level protocol in the artifact. Selective-split comparisons are reported in the appendix and illustrate how the audited target affects the supported decision.

Thresholds and metrics. Each task uses pre-specified preservation and localization thresholds recorded in the audited run packets. The main certificates use the bounded backend score `prob_diff` in $[-1, 1]$ and the signed transformed audit loss `prob_diff_delta` in $[-2, 2]$. Preservation is assessed with one-sided upper confidence bounds, localization with one-sided lower confidence bounds, and power through empirical-Bernstein minimum detectable effects. Stability is summarized by mean pairwise Jaccard overlap and core/fringe decomposition. Splits are formed at the level of independent task-instance rows, before construction of derived patching evaluations.

5 Results

5.1 Selector-level evidence separates into regimes

Table 2 reports four selector-level regimes. For true selector-level cross-fit, the size column gives the median foldwise selected-circuit size rather than a single fixed-circuit size.

Case	$ C_k $ med. [min,max]	$\hat{R}_{\text{pres}}/\text{UCB}$	Pres.	$\hat{R}_{\text{loc}}/\text{LCB}$	Loc.	Robust. (P/L)	Power	Jaccard
IOI	49 [48,52]	0.721 / 0.802	×	0.873 / 0.790	✓	×/ ✓	✓	0.942
ARC-Easy	77 [70,84]	0.004 / 0.091	✓	0.004 / −0.082	×	✓/ ×	✓	0.821
Arithmetic	65 [50,72]	0.980 / 1.456	×	1.110 / 0.623	✓	×/ ×	×	0.647
InterpBench / EAP-IG	59 [57,62]	0.161 / 0.251	×	0.440 / 0.325	✓	×/ ✓	✓	0.791

Table 2: Main selector-level procedure audits. The size column reports the median foldwise selected-circuit size. Preservation and localization decisions are thresholded separately from the adequate-power flag, and robustness is decomposed into preservation robustness (P) and localization robustness (L) over the declared intervention family $\{\text{counterfactual}, \text{mean}\}$.

Object	$ C $	Behav. (P/L)	Precision	F1
Ground truth	7	×/ ×	1.000	1.000
Full graph	1108	✓/ ×	0.006	0.013
Empty graph	0	×/ ×	0.000	0.000
EAP / true selector	38	×/ ✓	0.026	0.044
EAP-IG / true selector	53	×/ ✓	0.019	0.033

Table 3: InterpBench reports behavioral support alongside edge-level structural recovery. The selected EAP and EAP-IG circuits support localization under the audited intervention while remaining weak recoveries of the ground-truth edge set.

Localization-dominant regime. On GPT-2 small IOI, true selector-level cross-fit returns foldwise circuits with median size 49 and range [48, 52], high selector stability (Jaccard 0.942), and adequate power (MDE 0.081). Localization is supported, with estimate 0.873 and one-sided lower bound 0.790, giving a margin of 0.490 above the threshold. Preservation falls short: the estimate is 0.721 and the one-sided upper bound is 0.802. Minimality fails under Holm-adjusted edgewise testing. The robustness vector follows the behavioral split: localization robustness passes over the declared intervention family $\{\text{counterfactual}, \text{mean}\}$, while preservation robustness fails. The resulting selector-level profile is localization without preservation. The same qualitative pattern appears on Gemma2 arithmetic subtraction, with weaker power.

Preservation-dominant regime. On Gemma2 ARC-Easy, the selector-level audit exhibits the converse profile. The foldwise circuits have median size 77 and range [70, 84]. Preservation is supported, with estimate 0.004 and upper confidence bound 0.091, giving a margin of 0.109 below the threshold; preservation robustness also passes. Localization fails, with estimate 0.004 and a lower confidence bound below zero. The resulting profile is preservation without localization: the selected circuit preserves the task metric, while the audit does not concentrate the behavior in the selected subgraph.

Minimality and stability. All four main selector-level cases fail familywise minimality. Selector stability ranges from very high on IOI to moderate on arithmetic. Stability measures concentration of the selector distribution, while minimality measures edgewise essentiality under the audit. The main cases therefore include selectors that repeatedly return similar circuits whose audited edge effects still contain a broad unsupported fringe.

Power. The arithmetic case illustrates the role of power in the certificate. Localization passes under the chosen threshold, while the empirical-Bernstein MDE is much larger than in IOI or ARC-Easy. The certificate records the positive localization decision together with its lower inferential resolution.

5.2 Behavioral support separates from structural recovery

InterpBench allows behavioral certificates to be compared against known ground-truth circuits. In this setting, the ground-truth circuit has perfect structural recovery by construction, whereas its behavioral audit depends on the audited intervention, score, thresholding convention, and graph granularity. The full graph preserves behavior but has low precision as a structural hypothesis. The selected EAP and EAP-IG circuits under true selector-level cross-fit localize behavior, although their structural




Figure 2: Preservation and localization across the main cases.




Figure 3: Behavioral support and structural recovery on InterpBench.

257 precision and F1 remain low. Behavioral support and structural recovery are therefore reported as
 258 separate evidential dimensions.

259 The MI results should be read with this separation in view. A selector can identify a behaviorally
 260 meaningful subgraph with limited recovery of the true edge-level mechanism, and a structurally exact
 261 circuit can behave differently under the behavioral objective induced by a particular intervention
 262 design. When the benchmark ground-truth circuit fails the behavioral audit, the result indicates a
 263 mismatch between the structural target and the audited behavioral estimand. This mismatch is part of
 264 the evidence structure captured by the certificate.

5.3 Protocol choice changes the supported claim

The appendix compares selective split with true selector-level cross-fit on shared tasks. On ARC-Easy, the selector-level protocol supports preservation where the selected-circuit audit did not; on both InterpBench methods, it supports localization where the selected-circuit audit did not; and on IOI, both protocols yield localization without preservation. These comparisons place the audit protocol inside the scientific claim, since changing the protocol changes the audited target.

5.4 Minimality profiles reveal descriptive high-effect tails

Although the Holm-adjusted familywise minimality decision fails on all main selector-level circuits, the minimality profiles reveal heavy-tailed edge-effect distributions. For IOI, the median edge effect is positive and the upper quartile is substantially larger than zero; the same qualitative shape appears on arithmetic and InterpBench. The appendix visualizes these profiles and the resulting picture is a descriptive high-effect tail embedded in a broad low-effect fringe, a structure that a single pass/fail minimality decision cannot capture.

6 Conclusion

Circuit discovery produces adaptive objects: subgraphs selected through task definitions, counterfactual constructions, scores, selectors, thresholds, and graph abstractions. The corresponding scientific unit of evidence is the mechanistic claim, which couples the selected circuit to the model, task, intervention design, audit protocol, audited property, and selection procedure under which it was obtained. A Circuit Certificate records that claim together with its uncertainty calculation, stability summary, minimality analysis, robustness vector, power calculation, and structural-recovery results when available.

The theory supplies the inferential discipline for this claim-level view. Selective split yields finite-sample one-sided certification for any measurable selector under independent audit data. True selector-level cross-fit gives valid foldwise out-of-fold audits for circuits selected without the audited fold and targets the behavior of the selection procedure across repeated rediscovery. Edgewise minimality can be audited with familywise error control, and power is reported in the same finite-sample regime as the behavioral bounds. Together, these components define circuit evaluation as a protocol-indexed statistical audit.

The experiments separate the evidential dimensions captured by the certificates. Preservation, localization, minimality, stability, robustness, power, and structural recovery vary across canonical and benchmarked settings. GPT-2 small IOI yields strong selector-level localization with high stability but unsupported preservation and minimality. Gemma2 ARC-Easy yields preservation without localization. InterpBench shows behavioral support coexisting with weak edge-level structural recovery. The certificate format reports these regimes as distinct properties of the claim, instead of compressing them into a single circuit-quality score.

Circuit Certificates provide an executable reporting standard for mechanistic interpretability by identifying the claim supported by a discovered circuit, attaching that claim to the protocol that produced the evidence, and reporting uncertainty in the vocabulary of the mechanistic property under audit. The resulting analysis reports auditable mechanistic evidence together with the selected subgraph.

References

- [1] Arthur Conmy, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/34e1dbe95d34d7ebaf99b9bcaeb5b2be-Abstract-Conference.html.
- [2] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt,

- 314 Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and
315 Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*,
316 2021. URL <https://transformer-circuits.pub/2021/framework/index.html>.
- 317 [3] Atticus Geiger, Duligur Ibeling, Amir Zur, Maheep Chaudhary, Sonakshi Chauhan, Jing Huang,
318 Aryaman Arora, Zhengxuan Wu, Noah Goodman, Christopher Potts, and Thomas Icard. Causal
319 abstraction: A theoretical foundation for mechanistic interpretability. *Journal of Machine*
320 *Learning Research*, 26(83):1–64, 2025. URL [https://jmlr.org/papers/v26/23-0058](https://jmlr.org/papers/v26/23-0058.html).
321 [html](https://jmlr.org/papers/v26/23-0058.html).
- 322 [4] Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. Localizing model
323 behavior with path patching. *arXiv preprint arXiv:2304.05969*, 2023. doi: 10.48550/arXiv.
324 2304.05969. URL <https://arxiv.org/abs/2304.05969>.
- 325 [5] Rohan Gupta, Iván Arcuschin, Thomas Kwa, and Adrià Garriga-Alonso. InterpBench:
326 Semi-synthetic transformers for evaluating mechanistic interpretability techniques. In
327 *Advances in Neural Information Processing Systems*, volume 37, pages 92922–92951, 2024.
328 doi: 10.52202/079017-2950. URL [https://proceedings.neurips.cc/paper_files/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/a8f7d43ae092d9a5295775eb17f3f4f7-Abstract-Datasets_and_Benchmarks_Track.html)
329 [paper/2024/hash/a8f7d43ae092d9a5295775eb17f3f4f7-Abstract-Datasets_and_](https://proceedings.neurips.cc/paper_files/paper/2024/hash/a8f7d43ae092d9a5295775eb17f3f4f7-Abstract-Datasets_and_Benchmarks_Track.html)
330 [Benchmarks_Track.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/a8f7d43ae092d9a5295775eb17f3f4f7-Abstract-Datasets_and_Benchmarks_Track.html).
- 331 [6] Michael Hanna, Sandro Pezzelle, and Yonatan Belinkov. Have faith in faithfulness: Going
332 beyond circuit overlap when finding model mechanisms. In *First Conference on Language*
333 *Modeling*, 2024. URL <https://openreview.net/forum?id=TZ0CCGDcuT>.
- 334 [7] Stefan Heimersheim and Neel Nanda. How to use and interpret activation patching. *arXiv*
335 *preprint arXiv:2404.15255*, 2024. doi: 10.48550/arXiv.2404.15255. URL <https://arxiv.org/abs/2404.15255>.
336 [org/abs/2404.15255](https://arxiv.org/abs/2404.15255).
- 337 [8] Maximilian Li and Lucas Janson. Optimal ablation for interpretability. In *Advances in*
338 *Neural Information Processing Systems*, volume 37, pages 109233–109282, 2024. doi:
339 10.52202/079017-3468. URL [https://proceedings.neurips.cc/paper_files/paper/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/c55e6792923cc16fd6ed5c3f672420a5-Abstract-Conference.html)
340 [2024/hash/c55e6792923cc16fd6ed5c3f672420a5-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/c55e6792923cc16fd6ed5c3f672420a5-Abstract-Conference.html).
- 341 [9] Nicolai Meinshausen and Peter Bühlmann. Stability selection. *Journal of the Royal Statistical*
342 *Society: Series B (Statistical Methodology)*, 72(4):417–473, 2010. doi: 10.1111/j.1467-9868.
343 2010.00740.x.
- 344 [10] Maxime Méloux, François Portet, and Maxime Peyrard. Mechanistic interpretability as statistical
345 estimation: A variance analysis. *arXiv preprint arXiv:2510.00845*, 2025. doi: 10.48550/arXiv.
346 2510.00845. URL <https://arxiv.org/abs/2510.00845>.
- 347 [11] Joseph Miller, Bilal Chughtai, and William Saunders. Transformer circuit evaluation metrics are
348 not robust. In *First Conference on Language Modeling*, 2024. URL [https://openreview](https://openreview.net/forum?id=zSf8PJyQb2).
349 [net/forum?id=zSf8PJyQb2](https://openreview.net/forum?id=zSf8PJyQb2).
- 350 [12] Aaron Mueller, Atticus Geiger, Sarah Wiegreffe, Dana Arad, Iván Arcuschin, Adam Belfki,
351 Yik Siu Chan, Jaden Fiotto-Kaufman, Tal Haklay, Michael Hanna, Jing Huang, Rohan
352 Gupta, Yaniv Nikankin, Hadas Orgad, Nikhil Prakash, Anja Reusch, Aruna Sankaranarayanan,
353 Shun Shao, Alessandro Stolfo, Martin Tutek, Amir Zur, David Bau, and Yonatan Belinkov. MIB:
354 A mechanistic interpretability benchmark. In *Proceedings of the 42nd International Conference*
355 *on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 45069–
356 45108. PMLR, 2025. URL <https://proceedings.mlr.press/v267/mueller25a.html>.
- 357 [13] Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter.
358 Zoom in: An introduction to circuits. *Distill*, 5(3), March 2020. doi: 10.23915/distill.00024.001.
359 URL <https://distill.pub/2020/circuits/zoom-in>.
- 360 [14] Claudia Shi, Nicolas Beltran-Velez, Achille Nazaret, Carolina Zheng, Adrià Garriga-
361 Alonso, Andrew Jesson, Maggie Makar, and David M. Blei. Hypothesis testing
362 the circuit hypothesis in LLMs. In *Advances in Neural Information Process-*
363 *ing Systems*, volume 37, pages 94539–94567, 2024. doi: 10.52202/079017-2997.
364 URL [https://proceedings.neurips.cc/paper_files/paper/2024/hash/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/abccb8a90b30d45b948360ba41f5a20f-Abstract-Conference.html)
365 [abccb8a90b30d45b948360ba41f5a20f-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/abccb8a90b30d45b948360ba41f5a20f-Abstract-Conference.html).

[15] Kevin Ro Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=NpsVSN6o4u1>.

A Proofs and formal details

A.1 Set-up

Let $Z = (X, \tilde{X}, Y)$ denote a generic task tuple, and let $w(c, Z)$ be a bounded measurable loss associated with a circuit c and a sampled tuple Z . Preservation and localization losses are special cases of this template. For a selected circuit C , the target risk is

$$R(C) = \mathbb{E}_{Z \sim \mathcal{D}}[w(C, Z)].$$

Under selective split, the selected circuit depends only on selector-visible data and is independent of the audit sample. Under true selector-level cross-fit, each foldwise selected circuit is independent of its own audit fold by construction.

A.2 Proof of Theorem 2

Proof. Condition on the selector-visible data S and therefore on the selected circuit $C = g(S)$. Since the audit sample is independent of S , the audit observations are i.i.d. from $\mathcal{D}$ conditional on S . The variables $W_i = w(C, Z_i)$ are therefore i.i.d. and bounded in $[a, b]$ conditional on S , with mean $R(C)$. Any valid one-sided confidence procedure for bounded i.i.d. means applies conditionally on S , including the empirical-Bernstein bound used in the artifact. Integrating over S gives the marginal coverage statement. $\square$

A.3 Proof of Proposition 3

Proof. Fix the fold partition and a fold k . The circuit C_k is selected without using the data in $\mathcal{I}_k$. Hence, for every $i \in \mathcal{I}_k$, the held-out tuple Z_i is independent of the circuit used to audit it. Conditional on C_k , the held-out losses $\{w(C_k, Z_i) : i \in \mathcal{I}_k\}$ are i.i.d. bounded variables with conditional mean

$$R_k = \frac{1}{|\mathcal{I}_k|} \sum_{i \in \mathcal{I}_k} \mathbb{E}[w(C_k, Z_i) \mid C_k].$$

Therefore any valid one-sided confidence procedure for bounded i.i.d. means applies on the held-out fold, including the empirical-Bernstein rule used in the artifact. $\square$

Remark on aggregation. The aggregated selector-level estimator remains a useful procedure-level summary of foldwise out-of-fold losses. However, because an observation in fold j may be used to select the circuit audited on another fold $k \neq j$, conditioning on all selected circuits does not in general restore independence across all aggregated terms for arbitrary adaptive selectors. We therefore report the aggregate as a procedure-level estimator without claiming an unconditional finite-sample concentration theorem in this setting.

A.4 Proof of Proposition 4

Proof. Condition on the selected circuit and on the data used to select it, leaving the edgewise audit data random. By assumption, each edgewise p -value is valid for its corresponding null hypothesis. Holm’s step-down procedure controls the familywise error rate at level α for any finite family of valid p -values. The conditional familywise guarantee therefore holds, and marginal control follows by averaging over the selection randomness. $\square$

A.5 Minimality profiles are monotone

Proposition 5. For any fixed circuit C , the function $\epsilon \mapsto \text{MinProf}_C(\epsilon)$ is monotone non-increasing.

406 *Proof.* If $\epsilon_1 \leq \epsilon_2$, then for every edge $e \in C$,

$$\mathbf{1}[\Delta_e(C; \eta) > \epsilon_2] \leq \mathbf{1}[\Delta_e(C; \eta) > \epsilon_1].$$

407 Summing over edges and dividing by $|C|$ gives the result. $\square$

408 A.6 Preservation and localization are logically distinct

409 **Proposition 6.** *There exist tasks, intervention designs, and circuits for which preservation holds and*
 410 *localization fails, and there exist tasks, intervention designs, and circuits for which localization holds*
 411 *and preservation fails.*

412 *Proof.* Normalize the clean-model score to 1. Consider a circuit whose selected-subgraph score is
 413 0.95 and whose complement score is 0.94. The preservation discrepancy is 0.05, and the localization
 414 discrepancy is 0.06. With thresholds $(0.1, 0.1)$, preservation holds and localization fails.

415 Conversely, consider a circuit whose selected-subgraph score is 0.05 and whose complement score
 416 is 0. The preservation discrepancy is 0.95, and the localization discrepancy is 1.0. With the same
 417 thresholds, localization holds and preservation fails. $\square$

418 A.7 Minimum detectable effects

419 Let n be the audit sample size, let $\hat{\sigma}^2$ be the audit variance estimate, and let $[a, b]$ be the range of the
 420 bounded loss. The artifact reports three minimum detectable effects:

$$\text{MDE}_{\text{normal}} = \frac{(z_{1-\alpha} + z_{\text{power}})\hat{\sigma}}{\sqrt{n}}, \quad \text{MDE}_{\text{Hoeffding}} = (b - a)\sqrt{\frac{\log(1/\alpha)}{2n}},$$

421

$$\text{MDE}_{\text{EB}} = \sqrt{\frac{2\hat{\sigma}^2 \log(3/\alpha)}{n}} + \frac{3(b - a) \log(3/\alpha)}{n}.$$

422 The main experiments use MDE_{EB} , aligning the power rule with the one-sided confidence bounds
 423 used for preservation and localization.

424 B Additional results

425 B.1 Protocol comparisons

426 Table 4 compares the audited protocols on shared tasks.

Case	Left protocol	Right protocol	Pres. (L→R)	Loc. (L→R)	$\Delta \hat{R}_{\text{pres}}$	$\Delta \hat{R}_{\text{loc}}$
IOI	selective split	true selector	×→×	✓→✓	0.167	0.022
InterpBench / EAP-IG	selective split	true selector	×→×	×→✓	-0.113	-0.056
InterpBench / EAP	selective split	true selector	×→×	×→✓	-0.077	-0.066
ARC-Easy	selective split	true selector	×→✓	×→×	-0.000	-0.000
Arithmetic	selective split	true selector	×→×	×→✓	0.062	0.085

Table 4: Protocol comparisons on shared tasks. The left-hand protocol audits one selected circuit, while the right-hand protocol audits the selector as a procedure. The transitions therefore reflect a change in audited target, not one protocol dominating a fixed-circuit claim.

427 B.2 Selector-level detailed values

428 Table 5 reports the selector-level values used for the threshold decisions.

429 B.3 InterpBench full suite

430 Table 6 gives the full InterpBench ground-truth suite.

Case	95% UCB_{pres}	m_{pres}	95% LCB_{loc}	m_{loc}	MDE	Jaccard
IOI / true selector	0.802	-0.502	0.790	0.490	0.081	0.942
InterpBench / EAP-IG / true selector	0.251	-0.051	0.325	0.125	0.091	0.791
ARC-Easy / true selector	0.091	0.109	-0.082	-0.282	0.087	0.821
Arithmetic subtraction / true selector	1.456	-1.256	0.623	0.423	0.476	0.647

Table 5: Detailed selector-level values. The margins are $m_{pres} = \tau_{pres} - UCB_{pres}$ and $m_{loc} = LCB_{loc} - \tau_{loc}$, so positive values indicate threshold clearance.

Object	Protocol	$ C $	Pres.	Loc.	Precision	Recall	F1
empty graph	fixed_holdout	0	×	×	0.000	0.000	0.000
full graph	fixed_holdout	1108	✓	×	0.006	1.000	0.013
ground truth	fixed_holdout	7	×	×	1.000	1.000	1.000
random same size eap seed0	fixed_holdout	50	×	×	0.020	0.143	0.035
random same size eap seed1	fixed_holdout	50	×	×	0.000	0.000	0.000
random same size eap seed2	fixed_holdout	50	×	×	0.000	0.000	0.000
random same size eapig seed0	fixed_holdout	59	×	×	0.017	0.143	0.030
random same size eapig seed1	fixed_holdout	59	×	×	0.000	0.000	0.000
random same size eapig seed2	fixed_holdout	59	×	×	0.000	0.000	0.000
eap selective split	selective_split	50	×	×	0.040	0.286	0.070
eap true crossfit selector	true_crossfit_selector	38	×	✓	0.026	0.143	0.044
eapig selective split	selective_split	59	×	×	0.034	0.286	0.061
eapig true crossfit selector	true_crossfit_selector	53	×	✓	0.019	0.143	0.033

Table 6: InterpBench ground-truth suite. The suite contrasts the ground-truth circuit, full and empty graphs, graph-size-matched random circuits, and circuits discovered by EAP and EAP-IG under selective-split and selector-level protocols.

431 B.4 Minimality profiles

432 Table 7 reports empirical minimality quantiles for the main cases.

Case	$q_{0.25}$	$q_{0.50}$	$q_{0.75}$	$q_{1.00}$
IOI / true selector	-0.0078	0.0029	0.0205	0.1189
ARC-Easy / true selector	-0.0000	-0.0000	0.0000	0.0010
Arithmetic / true selector	-0.0005	0.0014	0.0086	0.1758
InterpBench / EAP-IG / true selector	-0.0037	0.0106	0.0389	0.1861

Table 7: Empirical quantiles of edgewise minimality effects. The profiles describe a high-effect tail together with a wide low-effect bulk.




Figure 4: Minimality profiles for the main cases.




Figure 5: Selector stability for the main cases.

	Pres	Loc	Min	Stab	Loc-R	Pres-R	Pow	F1
GPT2 IOI true cross-fit	FAIL	PASS	FAIL	0.94	PASS	FAIL	PASS	n/a
InterpBench IOI EAP-IG true cross-fit	FAIL	PASS	FAIL	0.79	PASS	FAIL	PASS	0.033
InterpBench IOI EAP true cross-fit	FAIL	PASS	FAIL	0.56	PASS	FAIL	PASS	0.044
Gemma2 ARC-Easy true cross-fit	PASS	FAIL	FAIL	0.82	FAIL	PASS	PASS	n/a
Gemma2 Arithmetic true cross-fit	FAIL	PASS	FAIL	0.65	FAIL	FAIL	FAIL	n/a

Figure 6: Claim matrix for the main cases.

433 C Artifact, provenance, and reproducibility

434 The submission is supported by a consolidated artifact containing the code and archived run packets
435 used for the reported experiments.

436 **Protocol metadata.** Each certificate records the audit protocol, selector-independence status, and
437 identifiers for the discovery, calibration, and audit splits used in the run. These fields distinguish
438 fixed-circuit evidence from selector-level evidence throughout the archive.

439 **Execution.** The artifact README provides installation, a validation command, a minimal end-to-
440 end run, and direct paths to the archived run packets used by the reported cases.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction identify the paper’s three contributions: a claim-centered formalization, protocol-aware post-selection certificates, and empirical evidence that preservation, localization, minimality, stability, robustness, power, and structural recovery separate in practice. These contributions are developed in Sections 2–5.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The paper states the scope of the audited protocols, the dependence of certificates on graph, intervention, and task design, the role of audit sample size, and the distinction between behavioral support and structural recovery. The appendix summarizes additional practical limitations and artifact constraints.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The main text states the assumptions and theorem statements for selective-split validity, foldwise out-of-fold validity for selector-level cross-fit, and familywise-valid edgewise minimality. Appendix A provides the proofs and formal details, including the dependence conditions needed for stronger aggregate cross-fit concentration.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the audited tasks, selectors, protocols, thresholds, score ranges, and reported metrics. Appendix C describes the artifact. The supplementary bundle includes the code and archived run packets used to generate the reported results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The supplementary material contains the integrated implementation, and archived run packets used to produce the paper.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The paper specifies the audit protocols, audited tasks, primary selectors, calibration role, bounded score transformation, intervention family, and confidence rule used for statistical support. The appendix reports protocol comparisons and detailed values.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The main experimental claims are reported through one-sided confidence bounds for preservation and localization, familywise-valid minimality decisions, selector stability summaries, and power calculations aligned with the empirical-Bernstein rule used by the certificates.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The supplementary artifact includes executed run packets, protocol meta-data, and archived outputs. Resource-sensitive follow-ups, including optimal ablation, are separated from the main experimental spine.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is methodological and experimental. It uses existing research artifacts and benchmark settings, preserves anonymity in the submission package, and introduces no human-subject, privacy-sensitive, or high-risk deployment component.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper addresses evidential standards for mechanistic interpretability. Positive impacts include more reliable auditing of model mechanisms and more disciplined reporting of interpretability evidence. A plausible negative impact is that stronger circuit-evaluation methods could improve offensive capability analysis as well as defensive auditing; the paper discusses this risk at a high level.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no new high-risk generative model and no newly scraped dataset. The released materials consist of code and archived experimental results for auditing circuits in existing benchmarked settings.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: The paper credits the benchmark suites, selectors, and prior assets it builds on, including IOI, InterpBench, MIB, EAP, EAP-IG, and TransformerLens.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The submission introduces an integrated artifact for protocol-aware circuit certificates.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper involves no crowdsourcing or human-subject research.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

757 **15. Institutional review board (IRB) approvals or equivalent for research with human**
758 **subjects**

759 Question: Does the paper describe potential risks incurred by study participants, whether
760 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
761 approvals (or an equivalent approval/review based on the requirements of your country or
762 institution) were obtained?

763 Answer: [N/A]

764 Justification: The paper involves no human-subject research.

765 Guidelines:

766 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
767 with human subjects.

768 • Depending on the country in which research is conducted, IRB approval (or equivalent)
769 may be required for any human subjects research. If you obtained IRB approval, you
770 should clearly state this in the paper.

771 • We recognize that the procedures for this may vary significantly between institutions
772 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
773 guidelines for their institution.

774 • For initial submissions, do not include any information that would break anonymity (if
775 applicable), such as the institution conducting the review.

776 **16. Declaration of LLM usage**

777 Question: Does the paper describe the usage of LLMs if it is an important, original, or
778 non-standard component of the core methods in this research? Note that if the LLM is used
779 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
780 scientific rigor, or originality of the research, declaration is not required.

781 Answer: [N/A]

782 Justification: The scientific contribution concerns auditing mechanistic claims in existing
783 language-model selectors and benchmarks. LLMs are objects of study in the experiments.

784 Guidelines:

785 • The answer [N/A] means that the core method development in this research does not
786 involve LLMs as any important, original, or non-standard components.

787 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
788 be described.