
Score is Not Circuit: Evaluating Mechanistic Localization as a Score-to-Circuit Stack

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Mechanistic localization methods are often assessed through component-level im-
2 portance estimates, yet downstream claims rely on circuits induced from those
3 scores under particular selection rules, evaluation metrics, and compute budgets.
4 We argue that score quality and circuit quality are distinct objects, and we for-
5 malize mechanistic localization as a four-layer score-to-circuit stack consisting
6 of scoring, calibration, selection, and evaluation. We instantiate this view on
7 edge-level circuit localization and study how scorer corrections, uncertainty-aware
8 calibration, and metric-aware selectors interact across models and tasks in MIB.
9 Our analysis reveals three recurrent gaps: an estimation gap between approxi-
10 mate and exact edge scores, an interaction gap between edgewise importance and
11 multi-edge circuit faithfulness, and an objective gap between circuits favored by
12 performance-oriented and model-matching metrics. Across this stack, local scorer
13 improvements do not compose monotonically into better circuits: a scorer can
14 become more faithful to exact edge effects while yielding worse downstream cir-
15 cuits, and global uncertainty shrinkage can severely degrade circuit quality. These
16 findings support a simple conclusion: mechanistic localization should be evalu-
17 ated as a full score-to-circuit pipeline rather than as an isolated scoring problem.
18 We conclude by proposing compute-aware, metric-aware reporting standards for
19 future work on circuit discovery and localization.

1 Introduction

Mechanistic localization assigns importance scores to internal components, while downstream claims usually concern *circuits*: budgeted subgraphs produced by ranking, calibration, thresholding, or reranking those scores under explicit evaluation metrics and compute constraints [9, 2, 8, 6]. The gap between score fields and reported circuits now shapes empirical comparisons. Exact and approximate patching methods have improved [6, 4, 5], MIB moved evaluation from isolated case studies to explicit circuit metrics on fixed model–task cells [7], and the BlackboxNLP shared task showed that selection and regularization choices can materially change the circuits induced from the same score field [1]. These developments place the full score-to-circuit pipeline at the level of the scientific unit: the object that turns localization scores into evaluated circuit claims.

This paper develops an evaluation framework for mechanistic localization claims as *score-to-circuit pipelines*. We model each pipeline as a four-layer stack:

scorer → calibrator → selector → evaluator.

The guiding claim is:

A score field becomes a circuit claim only through the full stack.

Circuit quality is therefore a property of the scorer, calibrator, selector, and evaluator together.

We instantiate the score-to-circuit pipeline on edge-level circuit localization in MIB [7] by comparing exact edge patching, edge attribution patching, integrated-gradient scoring, and three edge-level nonlinear corrections inspired by AtP* [6]. The evaluation then compares annotation-only calibration with rewrite calibration and applies metric-aware selectors, including stable shortlists, shortlist-local lower-confidence reranking, and true PNR. Across the evaluated cells, agreement with exact edge effects has limited predictive power for downstream circuit quality, local scorer repairs compose nonmonotonically, global uncertainty rewrites can collapse strong raw rankings, and shortlist-based selectors often recover them by using uncertainty locally. Supported CPR winners and model-matching winners diverge repeatedly.

The resulting protocol formalizes mechanistic localization as a score-to-circuit stack, separates recurrent gaps in localization claims, and treats compute as a first-class reporting axis. It clarifies the metric geometry of circuit evaluation by distinguishing positive-performance quality, magnitude-mode area, and circuit–model distance. The accompanying MIB study shows that scorer quality, selector quality, and circuit quality are materially different objects. The released artifact makes the protocol auditable through proofs, executable witnesses, and mechanized checks.

Section 2 defines the stack, the metric geometry, and the structural results. Section 3 specifies the scorers, calibrators, selectors, and evaluation stages used in the empirical study. Section 4 presents the main findings. Section 5 distills the implications for future evaluation practice.

2 The Score-to-Circuit Stack

Definition 1 (Localization stack). *A localization stack is a quadruple $\mathfrak{L} = (S, C, A, \text{Eval})$ consisting of a scorer S , a calibrator C , a selector A , and an evaluator Eval .*

Given a model M , a clean/counterfactual dataset D , and an edge graph $G = (V, \mathcal{E})$, the scorer returns raw signed scores

$$S(M, D, G) = s \in \mathbb{R}^{|\mathcal{E}|}.$$

The calibrator maps raw scores to

$$C(s) = (\tilde{s}, \rho, \sigma),$$

where $\tilde{s}$ is a possibly rewritten score vector, ρ records stability information, and σ records uncertainty information. For an objective m , the selector constructs a family of budgeted circuits

$$A_m(\tilde{s}, \rho, \sigma) = \{C_k^m \subseteq \mathcal{E} : k \in K\},$$

where K is the benchmark budget grid. Finally, the evaluator returns circuit quality and cost,

$$\text{Eval}_m(\{C_k^m\}) = Q_m, \quad \text{Cost}(\mathfrak{L}) = (\text{wall time}, \text{VRAM}, \#F, \#B).$$

The localization stack

scores become circuit claims only after calibration, objective-specific selection, and evaluation




Figure 1: The localization stack. Mechanistic localization claims are produced by a pipeline that maps raw edge scores to budgeted circuits and then evaluates those circuits under objective-specific metrics and compute constraints.




Figure 2: Metric geometry for circuit evaluation. Circuit evaluation separates performance-oriented quality (CPR), the magnitude-mode area (AUC_{mag}), the benchmark distance (D_{MIB}), and the derived model-matching quality ($\text{MMQ} = 1 - D_{\text{MIB}}$). In the no-overshoot regime $f^{\text{mag}}(k) \leq 1$, $D_{\text{MIB}} = 1 - \text{AUC}_{\text{mag}}$ and therefore $\text{MMQ} = \text{AUC}_{\text{mag}}$.

2.1 Metric geometry

Positive-performance selection uses

$$u_e^+(s) = \max(s_e, 0),$$

and induces a faithfulness curve $f^+(k)$ whose integrated area is the circuit performance ratio

$$\text{CPR} = \int f^+(k) d\mu(k),$$

where μ is the normalized discrete measure on the benchmark budget grid. Magnitude-mode selection uses

$$u_e^{\text{mag}}(s) = |s_e|,$$

and induces a curve $f^{\text{mag}}(k)$ with auxiliary area

$$\text{AUC}_{\text{mag}} = \int f^{\text{mag}}(k) d\mu(k).$$

The benchmark distance is

$$D_{\text{MIB}} = \int |1 - f^{\text{mag}}(k)| d\mu(k),$$

with lower values preferred. For readability in winner tables and transfer tables, we also report the derived high-is-good quantity

$$\text{MMQ} = 1 - \text{D}_{\text{MIB}}.$$

Proposition 1 (Distance–area relation). *Assume that μ is normalized. Then*

$$\text{D}_{\text{MIB}} = 1 - \text{AUC}_{\text{mag}} + 2 \int (f^{\text{mag}}(k) - 1)_+ d\mu(k).$$

Hence, if $f^{\text{mag}}(k) \leq 1$ for all budgets, then $\text{D}_{\text{MIB}} = 1 - \text{AUC}_{\text{mag}}$ and therefore $\text{MMQ} = \text{AUC}_{\text{mag}}$. The correction term is exactly twice the total overshoot above 1.

CPR and MMQ define different geometric targets. Even within a fixed scorer family, their optima can occur at different stacks.

2.2 Four gaps

The stack induces four analytically distinct gaps. The *estimation gap* measures disagreement between approximate and exact edge scores on tractable cells. The *rank-distortion gap* measures how much a calibrator changes the raw ordering. The *interaction gap* measures the discrepancy between additive edgewise utility and realized multi-edge circuit faithfulness. The *objective gap* measures the disagreement between stacks favored by CPR and by model-matching quality. Compute is tracked throughout as a reporting axis because stack quality is meaningful only for materialized cells.

Two structural results organize the rest of the paper.

Proposition 2 (Invariance under isotonic calibration). *Let A_k be a deterministic top- k selector with fixed tie-breaking, and let ϕ be strictly increasing. Then*

$$A_k(\{\phi(u_e)\}) = A_k(\{u_e\})$$

for every budget k .

Rank-preserving calibration is separated from calibration procedures that rewrite the selection order. Annotation calibration preserves the raw ranking while allowing metadata-aware selectors to use ρ or σ near the selection boundary. Rewrite calibration can change the selected circuit through the induced utility order, tie structure, or selector inputs.

Proposition 3 (Zero interaction gap under exact additivity). *Assume there exists a non-negative edge utility system w^m such that every circuit’s faithfulness equals its normalized additive mass under w^m . Then the additive top- k proxy matches the realized faithfulness curve exactly, the interaction gap is zero, and the curve is monotone non-decreasing.*

Observed interaction gaps indicate limits of the additive score-to-circuit approximation. In practice, they may combine true multi-edge interaction, scorer estimation error, and selector–objective mismatch. The relevant evaluation unit is therefore the full stack.

3 Protocol and Experimental Setting

We instantiate the stack on edge-level circuit localization in MIB [7]. The supported matrix contains seven circuit-localization cells with both CPR and model-matching evaluation, together with InterpBench for edge-level AUROC [3].

3.1 Concrete stack instantiations

The scorer family consists of EAP, EAP-IG, exact edge patching on tractable cells, and three local nonlinear corrections inspired by AtP* [6]. In concise form,

$$s_e^{\text{EAP-QK}} = \langle \Delta a_e, \nabla_a L \rangle, \quad s_e^{\text{EAP-MLP}} = \langle \Delta \text{MLP}_e, \nabla_{\text{MLP}} L \rangle,$$

and EAP* applies both local corrections where the architecture supports them. Calibration is split into two classes:

- **annotation-only:** $\tilde{s}_e = s_e$, with (ρ_e, σ_e) exposed to the selector;

Item	What is reported
Cells	model, task, split/stage
Scorers	EAP, EAP-IG, Exact, EAP-QK, EAP-MLP, EAP*
Calibration	annotation-only vs. rewrite
Selectors	raw top- k , stable shortlist, LCB shortlist, true PNR
Hyperparameters	shortlist factor α , threshold τ , penalty λ , mixture p
Metrics	CPR $\uparrow$, AUC _{mag} $\uparrow$, D _{MIB} $\downarrow$, MMQ $\uparrow$, AUROC $\uparrow$
Compute	wall time, VRAM, bounded-probe status

Table 1: Reporting schema for score-to-circuit evaluation. Rows are stratified by scorer-level runs, selector development runs, frozen validation, and public-test replay.

Cell	Stage	CPR $\uparrow$ winner	MMQ $\uparrow$ winner
GPT-2 + IOI	supported val.	EAP-IG	EAP*
Gemma-2 + MCQA	supported val.	EAP-MLP	EAP-MLP
Gemma-2 + ARC-E	supported val.	EAP-MLP	EAP
Llama-3.1 + MCQA	supported val.	EAP-MLP	EAP*
Llama-3.1 + ARC-E	supported val.	EAP-IG	EAP
Qwen2.5 + IOI	supported val.	EAP-IG	EAP*
Qwen2.5 + MCQA	supported val.	EAP-IG	EAP-IG
InterpBench + IOI	supported val.	AUROC winner: EAP-IG	

Table 2: Metric winners on the supported validation matrix under the default scorer-level selector regime. The top stack for CPR and the top stack for MMQ differ in five of the seven circuit cells, showing that method rankings are objective-dependent.

- **rewrite**: the score field is modified, for example

$$\tilde{s}_e = \text{sign}(s_e) \max(|s_e| \rho_e - \lambda \sigma_e, 0).$$

- Selectors include signed top- k , magnitude top- k , true PNR, stable shortlists, and shortlist-local lower-confidence reranking. Stable shortlist methods protect the raw top- αk region and apply stability or uncertainty information inside that shortlist. LCB-within-shortlist uses

$$r_e^m = u_e^m - \lambda \sigma_e$$

- within the protected candidate set.

Compute accounting. Compute is part of the evaluation target. A method–cell pair is classified as *materialized* when it yields score files and downstream circuit evaluation under the declared budget. The remaining attempted pairs are classified as bounded probes or unsupported cells. Reported wall-clock times are end-to-end stack times for the corresponding metric evaluation.

Exact cells and frozen transfer. Exact edge patching is restricted to tractable and scientifically interpretable cells, namely GPT-2+IOI and InterpBench+IOI. Selector hyperparameters are chosen on a bounded development grid. A small number of stacks are then frozen and replayed on held-out validation and public test.

3.2 Selected comparisons

Cell	Frozen stack	Objective	Split	Quality	Time (s)
GPT-2 + IOI	EAP-MLP + stable shortlist	CPR $\uparrow$	val.	1.339	680
GPT-2 + IOI	same	CPR $\uparrow$	public	1.283	—
Qwen2.5 + IOI	EAP* + LCB shortlist	CPR $\uparrow$	val.	1.607	1589
Qwen2.5 + IOI	same	CPR $\uparrow$	public	1.604	—
Qwen2.5 + IOI	EAP* + LCB shortlist	MMQ $\uparrow$	val.	0.968	1579
Qwen2.5 + IOI	same	MMQ $\uparrow$	public	0.969	—
Qwen2.5 + MCQA	EAP* + stable shortlist	MMQ $\uparrow$	val.	0.922	193
Qwen2.5 + MCQA	same	MMQ $\uparrow$	public	0.919	—
Gemma-2 + MCQA	EAP-MLP + true PNR	CPR $\uparrow$	val.	1.696	274
Gemma-2 + MCQA	same	CPR $\uparrow$	public	1.720	—

Table 3: Selected frozen validation rows and public-test replay. Rows are reported when split, selector, and objective are unambiguous.

4 Results

4.1 Exact agreement and circuit quality diverge

We first examine the exact-tractable cells. On GPT-2+IOI, EAP-IG improves signed Pearson agreement with exact edge effects relative to EAP (0.931 vs. 0.814), while the strongest scorer-level CPR in that family comes from EAP-MLP rather than EAP-IG (2.177 vs. 2.087). On InterpBench, EAP-IG delivers the best AUROC, while EAP* collapses despite combining two local corrections. Exact one-edge fidelity and induced circuit quality therefore produce different empirical rankings.

The scorer ablation also clarifies the role of local nonlinear corrections. EAP-QK functions as a negative control: it provides little downstream rescue and degrades sharply on InterpBench. EAP-MLP is the strongest scorer-level rescue in this campaign, improving CPR over EAP on GPT-2+IOI and Qwen2.5+MCQA while remaining competitive on model matching. EAP* is useful but mixed: the combined scorer composes its local corrections nonmonotonically.

4.2 Interaction and objective gaps are structural

To evaluate the interaction gap, we compare complete stacks rather than isolated edge scores. Under exact additive composition, top- k utility mass would predict the realized faithfulness curve and the gap would vanish by Proposition 3. The empirical curves retain substantial residual gaps, indicating that a simple additive score-to-circuit approximation is insufficient. In practice, the residual may combine true multi-edge interaction, scorer estimation error, and selector-objective mismatch, so edgewise score quality alone leaves out information needed to predict circuit quality.

For CPR and MMQ, Table 2 shows different top stacks in five of the seven supported circuit cells. The disagreement is especially sharp on GPT-2+IOI and Qwen2.5+IOI, where EAP-IG is strongest for CPR but EAP* yields the strongest model-matching circuits. Separating performance-oriented and model-matching objectives changes the selected stack.

4.3 Global uncertainty rewrites can collapse strong scorers

Annotation and rewrite calibration induce different stack behavior. Annotation-only calibration preserves the raw score order and exposes metadata for boundary decisions. Rewrite calibration globally modifies the score field. The most dramatic case is Qwen2.5+IOI with EAP*. The raw-selector stack achieves CPR = 1.607 and MMQ = 0.968 on frozen validation and transfers almost unchanged to public test (1.604, 0.969). The corresponding rewrite collapses to CPR $\approx$ 0.25 and MMQ $\approx$ 0.25 on validation. Same-slice controls isolate the rewrite stage as the driver of the collapse.

This failure motivates a safer selector design: use uncertainty inside a protected candidate set. Stable shortlists and shortlist-local lower-confidence reranking recover several of the strongest cells because they preserve the raw support and apply uncertainty near the selection boundary. In several top rows, this selector-level decision changes circuit quality more than marginal scorer improvements.




Figure 3: Exact-score fidelity versus downstream circuit quality on tractable cells. Agreement with exact edge effects and induced circuit quality yield different empirical rank orders.




Figure 4: Representative interaction-gap cases. Each panel compares an additive edgewise proxy with the realized circuit faithfulness curve. The gaps are large, non-monotone, and budget-dependent, showing where additive score-to-circuit translation breaks down.

159 4.4 Compute changes the scientific recommendation

160 Stack comparisons depend on the compute regime. On the supported matrix, exact patching re-
 161 mains indispensable on tractable cells, while runtime and memory restrict its scope. EAP-IG often
 162 provides the strongest raw CPR, but selector design and frozen replay can change its practical
 163 advantage. Concrete examples illustrate the tradeoff. On InterpBench+IOI, EAP-IG improves AU-
 164 ROC over EAP while also being faster on the supported-matrix rerun (0.747 in 48.9s versus 0.726
 165 in 56.9s). On GPT-2+IOI, EAP-MLP reaches a stronger scorer-level CPR than EAP-IG at com-
 166 parable end-to-end cost (2.177 in 1313s versus 2.087 in 1505s). On Qwen2.5+MCQA, by contrast,
 167 EAP-IG retains an advantage in both CPR and MMQ while being slightly faster than the nonlinear
 168 variants.

169 Across the evaluated cells, score-to-circuit behavior depends on the full resource-bounded stack.
 170 Raw scorers paired with robust selectors are strongest in some settings, different scorers are favored




Figure 5: Calibration collapse and local rescue. Global rewrite can collapse strong raw scorers, whereas selectors that spend uncertainty within a protected shortlist preserve or recover several of the strongest cases. The sharpest collapse appears for Qwen2.5+IOI with EAP*.

by different objectives in others, and better agreement with exact edge effects can still yield weaker circuits. These patterns are stack-level properties and are not captured by scorer-only summaries.

5 Discussion: From Scorers to Stack Evaluation

The empirical study supports a change in reporting practice for mechanistic localization.

Stack-level claims. Local scorer improvements can compose nonmonotonically into downstream circuits. This is clearest for EAP-QK and for the mixed behavior of EAP*. Once circuits are selected, evaluated under metric-specific objectives, and constrained by compute, scorer-level summaries provide only one layer of the localization claim.

Exact edge fidelity as a causal reference. Exact patching provides the right one-edge causal reference on tractable cells. Circuit quality also depends on objective choice, selector design, and multi-edge composition. As a result, stronger agreement with exact edge effects can coexist with weaker induced circuits.

Calibration type. Annotation and rewrite are different scientific operations. Annotation-only calibration preserves the raw score field and exposes metadata for selection. Rewrite calibration globally changes the candidate order and can remove the useful support of a strong raw ranking. Localization papers should state whether uncertainty is used as annotation, as a local shortlist reranker, or as a global rewrite.

Metric orientation. This paper distinguishes CPR $\uparrow$, $\text{AUC}_{\text{mag}} \uparrow$, $D_{\text{MIB}} \downarrow$, and MMQ $\uparrow$. These quantities answer different geometric questions. In particular, model-matching quality and raw magnitude-mode area diverge once overshoot is admitted. Reporting them separately makes the objective gap visible.

Compute accounting. Claims about localization quality depend on what was materialized. Reporting should include the scorer, calibrator type, selector, objective, split/stage, wall time, and any bounded-probe or unsupported status. These fields identify the scientific object being evaluated.

These findings suggest the following reporting discipline:

1. Report the full stack (S, C, A, Eval) together with the scorer.
2. Distinguish annotation-only calibration from rewrite calibration.
3. Separate CPR, AUC_{mag} , D_{MIB} , and MMQ explicitly.
4. Use exact edge patching as a causal reference on tractable cells.
5. Report compute and admissibility together with quality.

The companion artifact supports this discipline as an auditability device. It includes full prose proofs, executable theorem witnesses, and mechanized checks for the central structural invariants.

6 Conclusion

Mechanistic localization claims reach empirical use as circuit claims: budgeted subgraphs selected from score fields under metric and compute constraints. We formalize this evaluation unit as a score-to-circuit stack, separating scorer, calibrator, selector, and evaluator, and distinguishing performance-oriented quality from model-matching quality. Under this formulation, localization claims can change through score estimation, rank transformation, multi-edge composition, objective choice, and materialized compute.

In the MIB-grounded study, these layers change the scientific conclusion. Exact-score agreement and downstream circuit quality rank methods differently; CPR and MMQ winners diverge in five of the seven supported circuit cells; local nonlinear corrections compose nonmonotonically; and global uncertainty rewrite can collapse a strong raw stack, while shortlist-based selectors can preserve and transfer useful support. Compute further determines which stacks are admissible for a given model-task cell.

Mechanistic localization papers should therefore report the full stack (S, C, A, Eval) , together with the calibration type, selector, objective, split or stage, quality metric, and compute status. This reporting unit makes localization claims comparable, auditable, and faithful to the object ultimately being evaluated: circuits induced from scores under declared constraints.

References

- [1] Dana Arad, Yonatan Belinkov, Hanjie Chen, Najoung Kim, Hosein Mohebbi, Aaron Mueller, Gabriele Sarti, and Martin Tutek. Findings of the BlackboxNLP 2025 shared task: Localizing circuits and causal variables in language models. In *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 543–552, Suzhou, China, November 2025. Association for Computational Linguistics. ISBN 979-8-89176-346-3. doi: 10.18653/v1/2025.blackboxnlp-1.32. URL <https://aclanthology.org/2025.blackboxnlp-1.32/>.
- [2] Arthur Conmy, Augustine Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems*, volume 36, pages 16318–16352, 2023. doi: 10.52202/075280-0719. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/34e1dbe95d34d7ebaf99b9bcaeb5b2be-Abstract-Conference.html.
- [3] Rohan Gupta, Iván Arcuschin, Thomas Kwa, and Adrià Garriga-Alonso. InterPBench: Semi-synthetic transformers for evaluating mechanistic interpretability techniques. In *Advances in Neural Information Processing Systems*, volume 37, pages 92922–92951, 2024. doi: 10.52202/079017-2950. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/a8f7d43ae092d9a5295775eb17f3f4f7-Abstract-Datasets_and_Benchmarks_Track.html.
- [4] Michael Hanna, Sandro Pezzelle, and Yonatan Belinkov. Have faith in faithfulness: Going beyond circuit overlap when finding model mechanisms. In *First Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=TZ0CCGDcuT>.
- [5] Farnoush Rezaei Jafari, Oliver Eberle, Ashkan Khakzar, and Neel Nanda. RelP: Faithful and efficient circuit discovery in language models via relevance patching. *arXiv preprint arXiv:2508.21258*, 2025. doi:

- 10.48550/arXiv.2508.21258. URL <https://arxiv.org/abs/2508.21258>. NeurIPS 2025 Mechanistic Interpretability Workshop.
- [6] János Kramár, Tom Lieberum, Rohin Shah, and Neel Nanda. AtP*: An efficient and scalable method for localizing LLM behaviour to components. *arXiv preprint arXiv:2403.00745*, 2024. doi: 10.48550/arXiv.2403.00745. URL <https://arxiv.org/abs/2403.00745>.
- [7] Aaron Mueller, Atticus Geiger, Sarah Wiegrefe, Dana Arad, Iván Arcuschin, Adam Belfki, Yik Siu Chan, Jaden Fried Fiotto-Kaufman, Tal Haklay, Michael Hanna, Jing Huang, Rohan Gupta, Yaniv Nikankin, Hadas Orgad, Nikhil Prakash, Anja Reusch, Aruna Sankaranarayanan, Shun Shao, Alessandro Stolfo, Martin Tutek, Amir Zur, David Bau, and Yonatan Belinkov. MIB: A mechanistic interpretability benchmark. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 45069–45108. PMLR, 2025. URL <https://proceedings.mlr.press/v267/mueller25a.html>.
- [8] Aaquib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 407–416, Miami, Florida, US, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.blackboxnlp-1.25. URL <https://aclanthology.org/2024.blackboxnlp-1.25/>.
- [9] Kevin Ro Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=NpsVSN6o4u1>.

A Proofs of the structural results

This appendix records the mathematical proofs used in the main text. The companion artifact additionally contains executable witnesses and mechanized checks.

A.1 Distance–area relation

Assume that μ is normalized. Write $f(k) = f^{\text{mag}}(k)$ and decompose

$$|1 - f(k)| = \begin{cases} 1 - f(k), & f(k) \leq 1, \\ f(k) - 1, & f(k) > 1. \end{cases}$$

Equivalently,

$$|1 - f(k)| = 1 - f(k) + 2(f(k) - 1)_+.$$

Integrating both sides against μ gives

$$D_{\text{MIB}} = \int |1 - f(k)| d\mu(k) = 1 - \int f(k) d\mu(k) + 2 \int (f(k) - 1)_+ d\mu(k).$$

Since $\int f(k) d\mu(k) = \text{AUC}_{\text{mag}}$, we obtain

$$D_{\text{MIB}} = 1 - \text{AUC}_{\text{mag}} + 2 \int (f^{\text{mag}}(k) - 1)_+ d\mu(k).$$

If $f^{\text{mag}}(k) \leq 1$ for all budgets, the overshoot term vanishes and $D_{\text{MIB}} = 1 - \text{AUC}_{\text{mag}}$. Negative values below 0 do not by themselves break the identity; only overshoot above 1 contributes the correction term. $\square$

A.2 Isotonic calibration invariance

Let $u_{e_1}, \dots, u_{e_n}$ be utilities sorted in descending order with fixed tie-breaking. If ϕ is strictly increasing, then

$$u_{e_i} > u_{e_j} \iff \phi(u_{e_i}) > \phi(u_{e_j}),$$

and equality is preserved as well. Therefore the total order induced by the utilities is unchanged under ϕ . A deterministic top- k selector with fixed tie-breaking depends only on that order, so its selected set is identical before and after transformation for every k . $\square$

280 A.3 Zero interaction gap under exact additivity

281 Assume

$$f_m(C) = \frac{\sum_{e \in C} w_e^m}{\sum_{e \in \mathcal{E}} w_e^m}$$

282 for a non-negative edge utility system w^m . Let C_k be the set of the k highest-weight edges. By
 283 construction, C_k maximizes the numerator among all size- k circuits. Hence the additive top- k proxy
 284 and the realized faithfulness coincide exactly at budget k . Since this holds for every budget, the
 285 interaction gap is zero. Monotonicity follows because adding one more edge with non-negative
 286 weight cannot decrease the numerator. $\square$

287 B Additional Results



Figure 6: Primary frontier view across representative cells.




Figure 7: Micro-validation scatter for local corrections on sampled edges.




Figure 8: Rank-distortion histograms induced by calibration.




Figure 9: Selector-at-scale behavior across representative scorer-cell pairs.




Figure 10: Global compute frontier across reporting metrics.

Cell	Method	Pearson _{signed}	Pearson _·	Sign agr.	Recall@1%	Recall@5%
GPT-2 + IOI	EAP	0.814	0.817	0.969	0.870	0.890
GPT-2 + IOI	EAP-IG	0.931	0.932	0.698	0.755	0.644
GPT-2 + IOI	EAP-QK	0.590	0.591	0.940	0.650	0.658
GPT-2 + IOI	EAP-MLP	0.817	0.817	0.969	0.889	0.907
GPT-2 + IOI	EAP*	0.590	0.589	0.940	0.656	0.661
InterpBench + IOI	EAP	0.557	0.566	0.887	0.333	0.593
InterpBench + IOI	EAP-IG	0.644	0.664	0.685	0.167	0.704
InterpBench + IOI	EAP-QK	0.251	0.221	0.745	0.167	0.038
InterpBench + IOI	EAP-MLP	0.614	0.623	0.889	0.500	0.593
InterpBench + IOI	EAP*	0.273	0.245	0.746	0.167	0.077

Table 4: Score fidelity to exact edge effects on tractable cells. Higher exact-fidelity statistics do not reliably imply better downstream circuits.

288 C Companion artifact

289 The companion artifact follows the paper’s score-to-circuit framing. It contains the experimental
290 stack used in the paper, exact-tractable probes, selector audits, compute records, theorem witnesses,
291 and mechanized checks.

292 D Limitations and broader impacts

293 **Limitations.** The study centers on edge-level circuit localization under the intervention proto-
294 col and metric suite implemented by MIB. Exact-score analyses are limited to cells where exact
295 one-edge patching is computationally tractable. Those exact references provide causal edge-level
296 anchors; circuit-level conclusions still depend on the full score-to-circuit stack. The broader con-
297 clusions are strongest for the supported model–task matrix studied here. We report point estimates
298 for most circuit metrics, with bootstrap statistics used primarily for calibration and selector design.
299 Public-test transfer is reported for a selected set of frozen stacks after development, rather than as
300 an exhaustive second campaign.

301 **Broader impacts and release considerations.** This work contributes methodological infrastruc-
302 ture for evaluating mechanistic localization. Better score-to-circuit evaluation can improve the re-
303 liability of interpretability claims, reduce overstatement from score-only analyses, and make com-
304 parisons between localization pipelines more scientifically meaningful. The same evaluation tools
305 may lower the cost of identifying behaviorally important subgraphs in capable models, which could
306 support targeted steering or intervention.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The paper’s central claim is methodological: mechanistic localization should be evaluated as a full score-to-circuit stack rather than as an isolated scoring problem. The abstract and introduction state exactly this contribution and the body of the paper supports it through formalization, controlled empirical comparisons, and stack-level analysis.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Appendix D includes a dedicated limitations paragraph describing the scope of the protocol, the role of tractable exact cells, and the boundaries of the study’s claims.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The main text states the assumptions used by the structural propositions, Appendix A contains the full proofs in prose, and the companion artifact includes executable witnesses together with a completed Lean development of the formal layer.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper specifies the benchmark setting, supported model-task matrix, scorer families, calibration and selector families, exact-tractable cells, and confirmatory transfer protocol.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The submission is accompanied by an anonymized companion artifact containing the code, source data, diagnostics, and formal layer required to reproduce the reported stack evaluations. Appendix C summarizes the artifact structure.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The protocol section states the benchmark setting, selector and calibration families, supported cells, exact-tractable anchors, and confirmatory replay design, while the appendices and artifact document the concrete data products and supplementary details used for the reported experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No] Justification: The paper reports point estimates for most circuit-level metrics and uses bootstrap statistics primarily for calibration and selector design, rather than as error bars, confidence intervals, or statistical significance tests for the main experimental claims.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.

469 • It is OK to report 1-sigma error bars, but one should state it. The authors should preferably
 470 report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality
 471 of errors is not verified.

472 • For asymmetric distributions, the authors should be careful not to show in tables or figures
 473 symmetric error bars that would yield results that are out of range (e.g., negative error rates).

474 • If error bars are reported in tables or plots, the authors should explain in the text how they
 475 were calculated and reference the corresponding figures or tables in the text.

476 **8. Experiments compute resources**

477 Question: For each experiment, does the paper provide sufficient information on the computer
 478 resources (type of compute workers, memory, time of execution) needed to reproduce the exper-
 479 iments?

480 Answer: [Yes]

481 Justification: Compute is treated as part of the evaluated object. The main paper defines the
 482 compute doctrine and reports end-to-end stack times for representative cells.

483 Guidelines:

484 • The answer [N/A] means that the paper does not include experiments.

485 • The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud
 486 provider, including relevant memory and storage.

487 • The paper should provide the amount of compute required for each of the individual experi-
 488 mental runs as well as estimate the total compute.

489 • The paper should disclose whether the full research project required more compute than the
 490 experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it
 491 into the paper).

492 **9. Code of ethics**

493 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS
 494 Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

495 Answer: [Yes]

496 Justification: The work studies public benchmarks and open-source tooling, does not involve
 497 human subjects, and focuses on transparent evaluation methodology, reproducibility, and re-
 498 sponsible reporting of limitations and release considerations.

499 Guidelines:

500 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

501 • If the authors answer [No], they should explain the special circumstances that require a devi-
 502 ation from the Code of Ethics.

503 • The authors should make sure to preserve anonymity (e.g., if there is a special consideration
 504 due to laws or regulations in their jurisdiction).

505 **10. Broader impacts**

506 Question: Does the paper discuss both potential positive societal impacts and negative societal
 507 impacts of the work performed?

508 Answer: [Yes]

509 Justification: Appendix D discusses both the positive methodological value of better evaluation
 510 practices and the possible negative use of improved localization tooling for targeted intervention
 511 or steering in capable models.

512 Guidelines:

513 • The answer [N/A] means that there is no societal impact of the work performed.

514 • If the authors answer [N/A] or [No], they should explain why their work has no societal
 515 impact or why the paper does not address societal impact.

516 • Examples of negative societal impacts include potential malicious or unintended uses (e.g.,
 517 disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deploy-
 518 ment of technologies that could make decisions that unfairly impact specific groups), privacy
 519 considerations, and security considerations.

520 • The conference expects that many papers will be foundational research and not tied to par-
 521 ticular applications, let alone deployments. However, if there is a direct path to any negative

522 applications, the authors should point it out. For example, it is legitimate to point out that
523 an improvement in the quality of generative models could be used to generate Deepfakes
524 for disinformation. On the other hand, it is not needed to point out that a generic algorithm
525 for optimizing neural networks could enable people to train models that generate Deepfakes
526 faster.

- 527 • The authors should consider possible harms that could arise when the technology is being
528 used as intended and functioning correctly, harms that could arise when the technology is
529 being used as intended but gives incorrect results, and harms following from (intentional or
530 unintentional) misuse of the technology.
- 531 • If there are negative societal impacts, the authors could also discuss possible mitigation strate-
532 gies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for
533 monitoring misuse, mechanisms to monitor how a system learns from feedback over time,
534 improving the efficiency and accessibility of ML).

535 11. Safeguards

536 Question: Does the paper describe safeguards that have been put in place for responsible release
537 of data or models that have a high risk for misuse (e.g., pre-trained language models, image
538 generators, or scraped datasets)?

539 Answer: [N/A]

540 Justification: The submission does not release new model weights, scraped datasets, or other
541 high-risk generative assets. The companion artifact focuses on evaluation code, source data
542 products, and formal witnesses.

543 Guidelines:

- 544 • The answer [N/A] means that the paper poses no such risks.
- 545 • Released models that have a high risk for misuse or dual-use should be released with necessary
546 safeguards to allow for controlled use of the model, for example by requiring that users adhere
547 to usage guidelines or restrictions to access the model or implementing safety filters.
- 548 • Datasets that have been scraped from the Internet could pose safety risks. The authors should
549 describe how they avoided releasing unsafe images.
- 550 • We recognize that providing effective safeguards is challenging, and many papers do not
551 require this, but we encourage authors to take this into account and make a best faith effort.

552 12. Licenses for existing assets

553 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the
554 paper, properly credited and are the license and terms of use explicitly mentioned and properly
555 respected?

556 Answer: [Yes]

557 Justification: The paper cites the public assets it builds on, including MIB and InterpBench.

558 Guidelines:

- 559 • The answer [N/A] means that the paper does not use existing assets.
- 560 • The authors should cite the original paper that produced the code package or dataset.
- 561 • The authors should state which version of the asset is used and, if possible, include a URL.
- 562 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 563 • For scraped data from a particular source (e.g., website), the copyright and terms of service
564 of that source should be provided.
- 565 • If assets are released, the license, copyright information, and terms of use in the package
566 should be provided. For popular datasets, paperswithcode.com/datasets has curated
567 licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- 568 • For existing datasets that are re-packaged, both the original license and the license of the
569 derived asset (if it has changed) should be provided.
- 570 • If this information is not available online, the authors are encouraged to reach out to the asset's
571 creators.

572 13. New assets

573 Question: Are new assets introduced in the paper well documented and is the documentation
574 provided alongside the assets?

575 Answer: [Yes]

576 Justification: The paper introduces an artifact containing new code, derived data products, theorem witnesses, and Lean sources.

577

578 Guidelines:

- 579 • The answer [N/A] means that the paper does not release new assets.
- 580 • Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- 581
- 582
- 583 • The paper should discuss whether and how consent was obtained from people whose asset is used.
- 584
- 585 • At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.
- 586

587 **14. Crowdsourcing and research with human subjects**

588 Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

589

590

591 Answer: [N/A]

592 Justification: The paper does not involve crowdsourcing or research with human subjects.

593

594 Guidelines:

- 594 • The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- 595
- 596 • Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- 597
- 598
- 599 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.
- 600

601 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**

602 Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

603

604

605

606 Answer: [N/A]

607 Justification: The paper does not involve crowdsourcing or research with human subjects.

608

609 Guidelines:

- 609 • The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- 610
- 611 • Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- 612
- 613
- 614 • We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- 615
- 616
- 617 • For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.
- 618

619 **16. Declaration of LLM usage**

620 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

621

622

623

624 Answer: [N/A]

625 Justification: LLMs are the objects being analyzed. The core contribution is an evaluation methodology for mechanistic localization pipelines.

626

627 Guidelines:

- 628 • The answer [N/A] means that the core method development in this research does not involve
629 LLMs as any important, original, or non-standard components.
- 630 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not be
631 described.